

tspDB

tspDB enables predictive query functionality in PostgreSQL by building an additional “prediction index” for a collection of time-series of interest.

A prediction index will allow for fast data retrieval, but for entries that are:

- 1- At a future time step (i.e. forecasting);
- 2- Missing/corrupted by noise (i.e. imputation)

Our paper here (<https://arxiv.org/abs/1903.07097>) provides more information about how tspDB works and its performance.

Installation

For installation instruction, go to the installation page here ([installation](#)).

Getting Started

The main functionalities of tspDB are enabling predictive queries, which are enabled via creating a prediction index on your time series table. The index is created via the function

`create_pindex()` . which you can use as follow:

```
select create_pindex('tablename','time_column_name',{'value_column_name'},'index_name')
```

To get you familiar with tspDB capabilities, we provided a testing function that will create a set of time series tables in your database. The test function will also create several prediction indices.

Run the function from your Postgres terminal

```
SELECT test_tspdb();
```

if you get at the last line

```
NOTICE: Pindices successfully created
```

then the test has passed. Now we can check the predictive indices the test has created through

```
SELECT * FROM list_pindices();
```

You will see the three predictive indices created by the test. Now let's create our own predictive index, which we will call 'pindex1' on the time series table `mixturets2`. The prediction index is created on the time column `time` and the value column `ts_7`:

```
SELECT create_pindex('mixturets2','time','{"ts_7"}','pindex1');
```

we can see our newly created index by running `list_pindices` again:

```
SELECT * FROM list_pindices();
```

Let's now use that prediction index to produce some *predictions*! let's for example predict at a time `t` that exists in the database. Effectively, we are *denoising* the existing observation or *imputing* a null observation. For example, at time 1, `ts_7` has a null value as you can see by running:

```
SELECT ts_7 FROM mixturets2 WHERE time = 1;
```

Let's *impute* this point by running:

```
SELECT * FROM predict('mixturets2','ts_7',1,'pindex1');
```

Which will return predictions as well as upper and lower bound for a 95% confidence interval. We can get a tighter bound with lower confidence by changing the confidence interval to, say 80%:

```
SELECT * FROM predict('mixturets2','ts_7',1,'pindex1', c=> 80);
```

The prediction index also supports forecasting queries using the same function. For example, you can forecast the value of column `ts_7` at time 100010, ten points ahead of what exists in the database by running:

```
SELECT * FROM predict('mixturets2','ts_7',100010,'pindex1');
```

In a similar fashion, you can execute range predictive queries using `predict()`. For example, we can *impute* the first hundred points of `ts_7` using:

```
SELECT * FROM predict('mixturets2','ts_7',0,100,'pindex1');
```

or *forecast* the next 10 points using:

```
SELECT * FROM predict('mixturets2','ts_7',100001,100010,'pindex1');
```

For more information about the functionalities of tspDB, visit the [API Reference \(API\)](#) page.

Examples

For further examples, check the python notebook examples here (https://github.com/AbdullahO/tspdb/blob/master/notebook_examples).

Demo

Take a look at our 2020 NeurIPS demo here

(<https://colab.research.google.com/drive/1XqiahxFoISDnTkS8O-E0uwglGNXcNQo0?usp=sharing>).

Contributing

Please visit our Github (<https://github.com/AbdullahO/tspdb/blob/master/CONTRIBUTING.md>) page for more information.

License

This work is licensed under the Apache 2.0 License.

Documentation built with MkDocs (<http://www.mkdocs.org/>).